

```

from sklearn.datasets import load_iris
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
import matplotlib.pyplot as plt
import pandas as pd

df = pd.read_csv("/kaggle/input/spaceship-titanic/train.csv")
df.head()

```

	PassengerId	HomePlanet	CryoSleep	Cabin	Destination	Age	VIP	\
0	0001_01	Europa	False	B/0/P	TRAPPIST-1e	39.0	False	
1	0002_01	Earth	False	F/0/S	TRAPPIST-1e	24.0	False	
2	0003_01	Europa	False	A/0/S	TRAPPIST-1e	58.0	True	
3	0003_02	Europa	False	A/0/S	TRAPPIST-1e	33.0	False	
4	0004_01	Earth	False	F/1/S	TRAPPIST-1e	16.0	False	

	RoomService	FoodCourt	ShoppingMall	Spa	VRDeck	
0	0.0	0.0	0.0	0.0	0.0	Maham
1	109.0	9.0	25.0	549.0	44.0	Juanna
2	43.0	3576.0	0.0	6715.0	49.0	Altark
3	0.0	1283.0	371.0	3329.0	193.0	Solam
4	303.0	70.0	151.0	565.0	2.0	Willy

	Transported
0	False
1	True
2	False
3	False
4	True

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8693 entries, 0 to 8692
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   PassengerId           8693 non-null   object
1   HomePlanet            8492 non-null   object
2   CryoSleep             8476 non-null   object
3   Cabin                 8494 non-null   object
4   Destination           8511 non-null   object
5   Age                   8514 non-null   float64

```

```

6  VIP      8490 non-null object
7  RoomService 8512 non-null float64
8  FoodCourt  8510 non-null float64
9  ShoppingMall 8485 non-null float64
10 Spa      8510 non-null float64
11 VRDeck    8505 non-null float64
12 Name      8493 non-null object
13 Transported 8693 non-null bool

```

```
dtypes: bool(1), float64(6), object(7)
```

```
memory usage: 891.5+ KB
```

```
df.head()
```

```

      HomePlanet CryoSleep Cabin Destination Age VIP
RoomService \
0      Europa      False B/0/P TRAPPIST-1e 39.0 False      0.0
1      Earth      False F/0/S TRAPPIST-1e 24.0 False     109.0
2      Europa      False A/0/S TRAPPIST-1e 58.0  True      43.0
3      Europa      False A/0/S TRAPPIST-1e 33.0 False      0.0
4      Earth      False F/1/S TRAPPIST-1e 16.0 False     303.0

```

```

      FoodCourt ShoppingMall Spa VRDeck Transported
0          0.0          0.0  0.0    0.0         False
1          9.0          25.0 549.0   44.0          True
2       3576.0          0.0 6715.0   49.0         False
3       1283.0         371.0 3329.0  193.0         False
4          70.0         151.0  565.0    2.0          True

```

```
X = df.drop(columns=['Transported'])
```

```
y = df['Transported']
```

```
y = y.to_numpy()
```

```
X.head()
```

```

      HomePlanet CryoSleep Cabin Age VIP RoomService FoodCourt \
0      Europa      False B/0/P 39.0 False          0.0          0.0
1      Earth      False F/0/S 24.0 False        109.0          9.0
2      Europa      False A/0/S 58.0  True         43.0       3576.0
3      Europa      False A/0/S 33.0 False          0.0       1283.0
4      Earth      False F/1/S 16.0 False        303.0          70.0

      ShoppingMall Spa VRDeck
0          0.0    0.0    0.0
1         25.0   549.0   44.0
2          0.0  6715.0   49.0

```

3	371.0	3329.0	193.0
4	151.0	565.0	2.0

```

from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import ColumnTransformer

X['HomePlanet'] = LabelEncoder().fit_transform(X['HomePlanet'])
X['CryoSleep'] = LabelEncoder().fit_transform(X['CryoSleep'])
X['FoodCourt'] =
StandardScaler().fit_transform(X['FoodCourt'].to_numpy().reshape(-
1,1))
X['ShoppingMall'] =
StandardScaler().fit_transform(X['ShoppingMall'].to_numpy().reshape(-
1,1))
X['Spa'] =
StandardScaler().fit_transform(X['Spa'].to_numpy().reshape(-1,1))
X['VRDeck'] =
StandardScaler().fit_transform(X['VRDeck'].to_numpy().reshape(-1,1))
X['VIP'] = LabelEncoder().fit_transform(X['VIP'])
X['Cabin'] = LabelEncoder().fit_transform(X['Cabin'])
X['Age'] =
StandardScaler().fit_transform(X['Age'].to_numpy().reshape(-1,1))
X['RoomService'] =
StandardScaler().fit_transform(X['RoomService'].to_numpy().reshape(-
1,1))
X.tail()

```

	HomePlanet	CryoSleep	Cabin	Age	VIP	RoomService
FoodCourt \						
8688	1	0	146	0.848924	1	-0.482039
2.989544						
8689	0	1	5280	-0.755179	0	-0.482039
0.490206						
8690	0	0	5285	-0.197230	0	-0.482039
0.490206						
8691	1	0	2131	0.221232	0	-0.482039
1.361151						
8692	1	0	2131	1.058155	0	-0.107410
2.796642						

	ShoppingMall	Spa	VRDeck
8688	-0.463857	2.249560	-0.254117
8689	-0.463857	-0.480047	-0.473852
8690	2.924236	-0.477118	-0.473852
8691	-0.463857	0.460086	2.831082
8692	-0.463857	-0.480047	-0.438219

```

X['HomePlanet'] = X['HomePlanet'].fillna(X['HomePlanet'].mean())
X['Cabin'] = X['Cabin'].fillna(X['Cabin'].mean())

```

```

X['VIP'] = X['VIP'].fillna(X['VIP'].mean())
X['RoomService'] = X['RoomService'].fillna(X['RoomService'].mean())
X['Age'] = X['Age'].fillna(X['Age'].mean())
X['CryoSleep'] = X['CryoSleep'].fillna(X['CryoSleep'].mean())
X['RoomService'] = X['RoomService'].fillna(X['RoomService'].mean())
X['FoodCourt'] = X['FoodCourt'].fillna(X['FoodCourt'].mean())
X['Spa'] = X['Spa'].fillna(X['Spa'].mean())
X['ShoppingMall'] = X['ShoppingMall'].fillna(X['ShoppingMall'].mean())
X['VRDeck'] = X['VRDeck'].fillna(X['VRDeck'].mean())

```

```

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.33, random_state=46)
model = LogisticRegression()
model.fit(X_train, y_train)
y_preds = model.predict(X_test)
y_preds = model.predict(X_test)
from sklearn.metrics import accuracy_score
accuracy_score(y_test, y_preds)

```

```

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge
(status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
0.7647263855001742
```

```
df2 = pd.read_csv("/kaggle/input/spaceship-titanic/test.csv")
```

```
pid = df2['PassengerId']
```

```
df2 = df2.drop(columns=["Destination", "Name", "PassengerId"])
```

```
df2['HomePlanet'] = LabelEncoder().fit_transform(df2['HomePlanet'])
```

```
df2['CryoSleep'] = LabelEncoder().fit_transform(df2['CryoSleep'])
```

```
df2['FoodCourt'] =
```

```
StandardScaler().fit_transform(df2['FoodCourt'].to_numpy().reshape(-1,1))
```

```
df2['ShoppingMall'] =
```

```
StandardScaler().fit_transform(df2['ShoppingMall'].to_numpy().reshape(-1,1))
```

```
df2['Spa'] =
```

```
StandardScaler().fit_transform(df2['Spa'].to_numpy().reshape(-1,1))
```

```

df2['VRDeck'] =
StandardScaler().fit_transform(df2['VRDeck'].to_numpy().reshape(-1,1))
df2['VIP'] = LabelEncoder().fit_transform(df2['VIP'])
df2['Cabin'] = LabelEncoder().fit_transform(df2['Cabin'])
df2['Age'] =
StandardScaler().fit_transform(df2['Age'].to_numpy().reshape(-1,1))
df2['RoomService'] =
StandardScaler().fit_transform(df2['RoomService'].to_numpy().reshape(-
1,1))

```

```
df2.tail()
```

```

/usr/local/lib/python3.10/dist-packages/pandas/io/formats/
format.py:1458: RuntimeWarning: invalid value encountered in greater
    has_large_values = (abs_vals > 1e6).any()
/usr/local/lib/python3.10/dist-packages/pandas/io/formats/format.py:14
59: RuntimeWarning: invalid value encountered in less
    has_small_values = ((abs_vals < 10 ** (-self.digits)) & (abs_vals >
0)).any()
/usr/local/lib/python3.10/dist-packages/pandas/io/formats/format.py:14
59: RuntimeWarning: invalid value encountered in greater
    has_small_values = ((abs_vals < 10 ** (-self.digits)) & (abs_vals >
0)).any()

```

	HomePlanet	CryoSleep	Cabin	Age	VIP	RoomService	
FoodCourt \							
4272	0	1	2679	0.376787	0	-0.361266	-
0.287719							
4273	0	0	3265	0.941066	0	-0.361266	
0.266790							
4274	2	1	602	NaN	0	-0.361266	-
0.287719							
4275	1	0	603	NaN	0	-0.361266	
1.466805							
4276	0	1	2680	1.011601	0	-0.361266	-
0.287719							

	ShoppingMall	Spa	VRDeck
4272	-0.316173	-0.271297	-0.249197
4273	-0.285857	-0.262344	-0.133705
4274	-0.316173	-0.271297	-0.249197
4275	-0.316173	-0.271297	0.170262
4276	-0.316173	-0.271297	-0.249197

```

df2['HomePlanet'] = df2['HomePlanet'].fillna(df2['HomePlanet'].mean())
df2['Cabin'] = df2['Cabin'].fillna(df2['Cabin'].mean())
df2['VIP'] = df2['VIP'].fillna(df2['VIP'].mean())
df2['RoomService'] =
df2['RoomService'].fillna(df2['RoomService'].mean())
df2['Age'] = df2['Age'].fillna(df2['Age'].mean())

```

```
df2['CryoSleep'] = df2['CryoSleep'].fillna(df2['CryoSleep'].mean())
df2['RoomService'] =
df2['RoomService'].fillna(df2['RoomService'].mean())
df2['FoodCourt'] = df2['FoodCourt'].fillna(df2['FoodCourt'].mean())
df2['Spa'] = df2['Spa'].fillna(df2['Spa'].mean())
df2['ShoppingMall'] =
df2['ShoppingMall'].fillna(df2['ShoppingMall'].mean())
df2['VRDeck'] = df2['VRDeck'].fillna(df2['VRDeck'].mean())

preds = model.predict(df2)
output = pd.DataFrame({'Passengerid':pid.to_numpy(),
                        'Transported':preds})

output.to_csv("/kaggle/working/submission.csv", index=False)
```